



PUBLIC / EXTERNAL

Executive Overview

Quenyx vOPS HUB — v1.0.0 · Document v2.1

Document Version	2.1
Software Version	v1.0.0
Applies To	Quenyx vOPS HUB v1.0.0
Classification	Public / External
Owner	Executive Office
Status	Released
Last Updated	2026-06-30
Document Type	Executive brief

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0; roadmap Phases 1–4 / Sprints 20–25; native monitoring; shared platform AI.
2.1	2026-06-30	Sprint 21 (Operations Intelligence) delivered: QynSight becomes a live AI consumer on the shared platform.

Table of Contents

- 1. Vision
- 2. vOPS HUB positioning
- 3. Problem statement
- 4. Why unified operations + compliance + AI matters
- 5. Current platform status (what exists today)
- 6. QynSight status
- 7. QynShield / QCIF status
- 8. Quenyx AI Platform status
- 9. What exists today vs roadmap
- 10. Defensible architecture (the moat)
- 11. Investor narrative
- 12. Customer value
- 13. Risks and limitations (stated plainly)
- 14. Official roadmap

01 — Executive Overview

Audience: Investors, board, executives, strategic partners. **Tone:** Non-technical, factual, investor-ready.

Status basis: v1.0.0 (Phases 1–3 completed; Phase 4 in progress, Sprints 20–21 delivered).

This document describes **what exists today** in the Quenyx codebase and what is **planned**. It contains **no fabricated customers, revenue, market sizing, benchmarks, or certifications**.

1. Vision

Quenyx is building the **vOPS HUB** — a *Virtual Operations Center* that unifies, in one platform, the three functions that enterprises today run in disconnected silos:

1. **Operations & observability** — knowing the real state of infrastructure and services.
2. **Compliance & governance** — proving alignment to regulatory frameworks with evidence.
3. **Applied AI** — turning that operational and compliance data into explainable guidance.

The thesis: operations, compliance, and AI are converging. The winning platform is the one that treats them as **one source of truth** with a **shared, governable AI layer** — rather than bolting a chatbot onto fragmented tools.

2. vOPS HUB positioning

Quenyx vOPS HUB is a **modular** platform. Each module ("Qyn...") addresses one operational domain, but all modules share the same workspace model, identity, entitlements, audit, and AI platform.

- Today the HUB ships **QynSight** (operations/monitoring) and **QynShield** (compliance / QCIF).
- The platform is architected so additional modules plug in without re-plumbing identity, billing, or AI.

3. Problem statement

Enterprises — particularly in the Saudi/GCC market Quenyx targets — face three compounding problems:

- **Operational blind spots.** Monitoring is fragmented across tools; teams lack one trustworthy picture of infrastructure health and capacity.
- **Compliance is manual and unprovable.** Mapping controls (e.g., NCA ECC), collecting evidence, finding gaps, and producing auditable recommendations is spreadsheet-driven and error-prone.
- **AI is risky to adopt.** Generic LLM tools hallucinate, cannot cite sources, and are unsafe for regulated environments where every claim must be traceable.

4. Why unified operations + compliance + AI matters

A single platform that already holds operational state **and** an official compliance corpus can do something point tools cannot: **answer governance questions with cited, deterministic evidence** and **ground AI in real platform data** instead of open-ended generation. The data moat (operational + corpus) and the governance model (provenance + citations + determinism) reinforce each other.

5. Current platform status (what exists today)

● Built and production-ready

- **QynSight v1.0** — real-time monitoring, infrastructure map, performance analytics, capacity planning, alert management, service checks, host/agent enrollment. **Feature-frozen.**
- **QCIF Compliance Engine (QynShield backend/API):** official-source corpus model, **NCA ECC-2:2024 Corpus Revision v1**, Knowledge Graph, Cross-Framework Mapping, Evidence Intelligence, Gap Assessment, Recommendation Engine — all UUID-based and provenance-tracked.
- **Deterministic Reasoning Engine** — rule-based reasoning that decides *what* to answer before any AI is involved.
- **Executive Demonstration Platform** — read-only dashboard, scorecard, timeline, explainability, and platform-metrics endpoints driven by **real engine data** (no fabricated numbers).
- **Quenyx AI Platform Foundation** — provider registry (mock + OpenAI), skill registry/router, module-adapter contract, and a dynamic capability catalog. **QynShield is the first live AI adapter; QynSight is now a live AI consumer via Operations Intelligence (Sprint 21).**
- **QynSight Operations Intelligence (Sprint 21)** — a Monitoring Copilot, alert explanation and investigation, deterministic root-cause analysis, incident timelines, and capacity/performance/infrastructure/service-health intelligence with evidence-based recommendations. It **reuses** the shared AI platform (no duplicated AI) and grounds every answer in **real** monitoring data.

● Built but feature-flagged (off by default)

- **Compliance Copilot v0** runs in **mock mode** by default; real-model mode is env-gated.
- **Retrieval / RAG Foundation** — provider-agnostic, **metadata-only**, with deterministic fallback; off by default and never indexes tenant evidence by default.

● Architecture-ready (contract only)

- Other HUB modules are **registered internally** but **hidden from the sidebar** by a frontend flag, with **reserved** AI-adapter contracts until each is built.

6. QynSight status

Production-ready, with the **core monitoring engine feature-frozen** at v1.0. With **Operations Intelligence (Sprint 21)** it now adds an explainable AI layer on top of that frozen engine — the platform's anchor module and the default-visible product in the UI.

7. QynShield / QCIF status

Backend and APIs are **built and exercised** (corpus, graph, mapping, evidence, gap, recommendation, copilot, executive). It is the **first production consumer of the Quenyx AI Platform**. UI surfacing beyond the executive/demo layer is part of the forward roadmap.

8. Quenyx AI Platform status

The AI runtime has been **extracted into a shared platform** that any module can consume. QynShield is live on it, and **QynSight is now live too** via Operations Intelligence (Sprint 21) — reusing the same providers, prompt orchestration, conversations, and audit with **no duplicated AI**. AI is **off by default, deterministic-first**, and **citation-enforced**.

9. What exists today vs roadmap

Area	Today	Roadmap
Monitoring	QynSight v1.0 (frozen) + Operations Intelligence (Sprint 21)	Deeper integrations
Compliance	QCIF + NCA ECC-2:2024 v1, full backend/API	More frameworks; richer UI
AI	Shared platform, mock default, deterministic + citations; live on QynShield and QynSight	Real-model GA, RAG GA, more module intelligence
Modules	QynSight visible; QynShield backend; others hidden	Progressive module enablement

10. Defensible architecture (the moat)

- **Data moat:** operational state **plus** an official, provenance-tracked compliance corpus.
- **Governance moat:** deterministic engines + "no citation, no answer" + UUID/provenance-first.
- **Platform moat:** a reusable AI layer with provider abstraction and module adapters, so new modules inherit safe AI without rebuilding it.

11. Investor narrative

Quenyx has completed a **disciplined foundation across Phases 1–3**: a production monitoring product (native QynSight engines), a provenance-grounded compliance engine, and a governable shared AI platform — built with explicit safety rails rather than demo shortcuts. **Phase 4 (Enterprise AI Platform)** is now underway, with the **Unified AI Workspace (Sprint 20, presented as "Quenyx AI")** and **Operations Intelligence (Sprint 21)** already delivered in v1.0.0 — the latter turning QynSight from a monitoring product into an operations *intelligence* product on the shared AI platform — converting this foundation into broader module coverage and customer-facing AI from a base that is already architecturally sound.

12. Customer value

- One platform for operations + compliance instead of a tool sprawl.
- **Auditable** compliance: every answer traces to an official source.
- **Safe AI**: explainable, cited, deterministic-first, off unless deliberately enabled.

13. Risks and limitations (stated plainly)

- QynShield's customer-facing **UI** is earlier-stage than its backend; today's surface is the executive/demo layer + APIs.
- **Real-model AI and RAG are feature-flagged**, not yet GA.
- Only **one framework (NCA ECC-2:2024)** is loaded today.
- **No production metrics, customers, or certifications are claimed** in this pack.
- Several modules exist as registry entries only and are **hidden in the UI**.

14. Official roadmap

Phase	Name	Status
Phase 1	Platform Foundation	✔ Completed
Phase 2	Operations Platform (QynSight)	✔ Completed
Phase 3	Compliance & Enterprise AI Foundation (QCIF Sprints 1–19, AI Platform Foundation)	✔ Completed
Phase 4	Enterprise AI Platform	● In progress

Phase 4 — Enterprise AI Platform

Sprint	Title	Status
Sprint 20	Unified AI Workspace ("Quenyx AI")	✔ Delivered in v1.0.0
Sprint 21	Operations Intelligence	✔ Delivered in v1.0.0
Sprint 22	Asset Intelligence + AI Adapter Platform	✔ Delivered in v1.0.0
Sprint 23	Automation & Response Intelligence	Planned
Sprint 24	Service, Notification & Cost Intelligence	Planned
Sprint 25	Enterprise Intelligence Platform	Planned

There is no roadmap content beyond Sprint 25 at this time. Hidden business modules (QynRun, QynKnow, QynNotify, QynReact, QynVA, QynSupport, QynBalance) already exist as registered platform modules and are **disabled in the navigation by a sidebar flag** until their production rollout.

Sprint 22 — Asset Intelligence + AI Adapter Platform. Quenyx AI became a true **platform**: module AI is now a discoverable **adapter** (registry + shared narrator) with **no per-module branching**, so every

future module plugs in the same way. **QynAsset** is the second production AI consumer (after QynSight), turning the **real discovered inventory** into explainable Asset Intelligence — discovery, CMDB questions, dependencies, hardware/capacity, lifecycle, and evidence-based recommendations. Facts with no data source (software licenses, warranty/EOL dates) are reported honestly as *not collected*, never fabricated. See Docs 22 (QynAsset Guide) and 23 (AI Adapter Developer Guide).