



PUBLIC / EXTERNAL

Executive Whitepaper

Quenyx vOPS HUB — v1.0.0 RC1 · Document v2.0

Document Version	2.0
Software Version	v1.0.0 RC1
Applies To	Quenyx vOPS HUB v1.0.0 RC1
Classification	Public / External
Owner	Office of the CTO
Status	Released
Last Updated	2026-06-29
Document Type	Whitepaper

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0 RC1; native monitoring; QynCore internal communication; shared platform AI.

Table of Contents

- | 1. The Virtual Operations Center concept
- | 2. Why operations, compliance, and AI converge
- | 3. Enterprise pain points
- | 4. Saudi / GCC compliance context
- | 5. AI governance model
- | 6. Why deterministic engines matter
- | 7. Why citations and provenance matter
- | 8. Quenyx roadmap (executive view)

04 — Executive Whitepaper

Audience: CIO, CTO, CISO, board. **Purpose:** Explain the concept and governance model behind Quenyx vOPS HUB without unsupported claims.

1. The Virtual Operations Center concept

Enterprises run operations, compliance, and (increasingly) AI as separate disciplines with separate tools and separate data. A **Virtual Operations Center (vOPS)** collapses these into one platform with **one source of truth**, one identity and entitlement model, one audit trail, and one **governable AI layer**. Quenyx vOPS HUB is that center: monitoring (QynSight) and compliance (QynShield/QCIF) sharing a workspace model and a common AI platform.

2. Why operations, compliance, and AI converge

- **Operations produces the facts** — what infrastructure exists and how it behaves.
- **Compliance interprets the facts** against regulatory frameworks and evidence.
- **AI explains and accelerates** — *if* it is grounded in those facts and constrained to cite them.

Run separately, each is weaker: monitoring lacks governance context, compliance lacks live operational truth, and AI lacks trustworthy grounding. Converged, they compound.

3. Enterprise pain points

- Fragmented monitoring → no single operational picture.
- Manual, spreadsheet-driven compliance → slow, error-prone, hard to audit.
- Generic LLM tools → hallucination, no provenance, unacceptable in regulated settings.
- Tool sprawl → duplicated identity, access, and audit surfaces, each a risk.

4. Saudi / GCC compliance context

Quenyx targets regulated enterprises in the Saudi/GCC market, where authorities such as the **National Cybersecurity Authority (NCA)** publish control frameworks. Quenyx ships the **NCA ECC-2:2024** framework as **Corpus Revision v1** (5 domains, 108 controls, 108 requirements), modelled from official source documents.

We make no claim of formal certification, endorsement, or accreditation by any authority. The corpus models official published controls; certification is the customer's responsibility.

5. AI governance model

Quenyx treats AI as a **renderer over an expert platform**, not as the expert:

- **Off by default.** No real model is called unless explicitly enabled (`AI_ENABLED`).
- **Deterministic first.** A rule-based Reasoning Engine decides *what* is true and *what* to answer before any model is involved.
- **Cited or silent.** The Copilot enforces citations — "no source, no answer."
- **Provider-agnostic.** Models are selected via a provider registry; no model is hardcoded; a mock provider allows safe operation with zero external calls.
- **Privacy-preserving.** Prompt logging and conversation persistence are **off by default**; tenant evidence is **never embedded** by default.

6. Why deterministic engines matter

Determinism makes outputs **reproducible and auditable**: the same inputs always yield the same reasoning and the same citations. In regulated environments, "the model felt confident" is not an acceptable basis for a control decision; "rule R-EVIDENCE-MISSING fired because requirement X has no evidence" is. Quenyx puts the deterministic engine in charge and uses AI only to phrase the result.

7. Why citations and provenance matter

Every compliance entity in Quenyx carries provenance back to an **official source document**, and every Copilot answer must cite the corpus. This:

- eliminates a whole class of AI hallucination risk,
- makes answers defensible to auditors,
- and lets customers verify, not just trust.

8. Quenyx roadmap (executive view)

- **Now (Phases 1–3 complete, v1.0.0 RC1):** QynSight production (native monitoring engines); QynShield backend/API; shared platform AI; deterministic reasoning; executive demonstration layer; Unified AI Workspace (Sprint 20) delivered.
- **Phase 4 — Enterprise AI Platform:** Sprint 20 Unified AI Workspace (delivered), Sprint 21 Operations Intelligence, Sprint 22 Asset & Knowledge Intelligence, Sprint 23 Automation & Response Intelligence, Sprint 24 Service/Notification/Cost Intelligence, Sprint 25 Enterprise Intelligence Platform. No roadmap beyond Sprint 25.

Bottom line for the board: Quenyx has built a governable foundation where AI is constrained by deterministic, provenance-backed engines — the prerequisite for safe AI adoption in regulated operations.