



INTERNAL

API Reference

Quenyx vOPS HUB — v1.0.0 · Document v2.2

Document Version	2.2
Software Version	v1.0.0
Applies To	Quenyx vOPS HUB v1.0.0
Classification	Internal
Owner	Platform Engineering
Status	Released
Last Updated	2026-06-30
Document Type	API reference

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0; includes Unified AI Workspace (Sprint 20) endpoints.
2.1	2026-06-30	Added QynSight Operations Intelligence (Sprint 21) endpoints under <code>`/api/qynsight/intelligence/*`</code> .
2.2	2026-06-30	Added the AI Adapter discovery API (<code>`/api/ai/adapters`</code> , <code>`/api/ai/actions`</code>) and QynAsset Asset Intelligence (<code>`/api/qynasset/intelligence/*`</code>) — Sprint 22.

Table of Contents

1. Authentication
2. Common headers
3. Error format
4. Workspace / project routes
5. QynSight APIs (summary) — project.module:qynsight
6. Compliance corpus APIs
7. AI context APIs (consumption contract, no AI execution)
8. Graph APIs
9. Mapping APIs
10. Evidence APIs
11. Gap APIs
12. Recommendation APIs
13. Copilot APIs — throttle compliance-copilot
14. Retrieval / RAG APIs
15. Executive APIs (read-only) — throttle compliance-executive
16. Platform AI capability endpoint
17. Quenyx AI (Unified AI Workspace — Sprint 20)
18. QynSight Operations Intelligence (Sprint 21)
19. QynAsset Asset Intelligence + AI Adapter Platform (Sprint 22)
- 19b. QynRun — Automation Platform (Sprint 23)
- 19c. QynReact — Incident Workspace (Sprint 23)
- 19d. QynKnow — Enterprise Knowledge Platform (Sprint 24)
- 19e. QynSupport — Service Desk + Ticket Intelligence (Sprint 24)
- 19f. QynNotify — Notification Intelligence (Sprint 24)
- 19g. Collaboration Platform (Sprint 24)
- 19h. QynVA — Enterprise AI Operator + Intelligence surfaces (Sprint 25)
- 19i. QynBalance — Enterprise Cost Intelligence (Sprint 25)
20. Notes on examples

08 — API Reference

Audience: Engineers, integrators. **Source:** Derived from `php artisan route:list` (261 routes) and the `routes/*.php` files at Sprint 19. **No endpoints are invented.** Regenerate this doc when routes change.

projects vs workspaces . Most tenant routes exist under **both** `/api/projects/{project}/...` and `/api/workspaces/{project}/...` — they are **aliases to the same controllers**. Below, paths are shown with `/workspaces/{project}` (the canonical form); swap `workspaces` → `projects` for the alias.

1. Authentication

- **Scheme:** Laravel **Sanctum** bearer tokens.
- **Obtain a token:** `POST /api/auth/login` (`POST /api/auth/register` to create an account).
- **Use it:** `Authorization: Bearer <token>` on every authenticated route.
- **Identity:** `GET /api/auth/me`; `logout POST /api/auth/logout` .
- **Public (no user auth):** `GET /api/health` , agent ingestion endpoints (token/secret auth).

2. Common headers

```
Authorization: Bearer <token>
Accept: application/json
Content-Type: application/json      # for POST/PUT/PATCH
```

3. Error format

Standard Laravel JSON errors:

```
{ "message": "Unauthenticated." }           // 401
{ "message": "This action is unauthorized." } // 403
{ "message": "Not Found" }                 // 404
{ "message": "The given data was invalid.", // 422
  "errors": { "field": ["validation message"] } }
{ "message": "Too Many Attempts." }       // 429 (throttled routes)
```

4. Workspace / project routes

Method	Path	Controller
GET/POST	/api/projects · /api/workspaces	ProjectController@index/sto
GET/PUT/DELETE	/api/workspaces/{project}	ProjectController@show/upda
GET	/api/workspaces/{project}/entitlements	ProjectSubscriptionControll
GET/PUT	/api/workspaces/{project}/subscription	ProjectSubscriptionControll
GET	/api/workspaces/{project}/modules · /modules/access	ProjectModuleController
PUT	/api/workspaces/{project}/modules/{moduleKey}/override	ProjectModuleOverrideContro
GET..DELETE	/api/workspaces/{project}/memberships...	ProjectMembershipController
GET	/api/workspaces/{project}/audit-logs	AuditLogController@index
GET..PUT	/api/workspaces/{project}/integrations...	ProjectIntegrationControlle
GET..POST	/api/workspaces/{project}/billing/...	BillingController
GET	/api/dashboard , /api/modules , /api/plans	platform meta

5. QynSight APIs (summary) — project.module:qynsight

Prefix: /api/workspaces/{project}/observe/...

- **Summary / instances:** summary , instances , instances/summary .
- **Services / checks:** services , service-definitions , run-checks .
- **Performance:** performance/metrics , real-time/metrics , real-time/system-info , real-time/thresholds .
- **Capacity:** capacity-planning , capacity-planning/export , capacity/metrics .
- **Alerts:** alerts/rules (GET/POST/PUT/DELETE/ toggle) , alerts/summary , alerts/history , alerts/events/{event}/acknowledge , alerts/channels , monitoring-profile (GET/PUT).
- **Infrastructure:** infrastructure/topology , infrastructure/connections , infrastructure/port-scans (GET/POST run).
- **Reports / data sources:** reports , reports/summary , data-sources , data-sources/summary .
- **Targets / hosts:** targets (GET/PUT), targets/validate , targets/{hostId}/port-scan .
- **Agents:** /api/workspaces/{project}/agents... (list, metadata, enrollment tokens, revoke, destroy).

Agent ingestion (token-auth, no user): GET /api/agents/download/{platform} , POST /api/agents/register , POST /api/agents/{agent}/heartbeat|metrics|inventory .

6. Compliance corpus APIs

Global (read): `/api/compliance/corpus/frameworks...` Workspace (entitled, `project.qynshield`): `/api/workspaces/{project}/compliance/corpus/frameworks...`

Path (suffix under <code>/corpus</code>)	Returns
<code>frameworks</code>	frameworks list
<code>frameworks/{frameworkKey}/releases</code>	releases
<code>frameworks/{frameworkKey}/releases/{releaseCode}/summary</code>	release summary
<code>.../domains</code> · <code>.../domains/{domainCode}</code>	domains
<code>.../controls/{controlCode}</code>	control detail
<code>.../search</code>	corpus search

7. AI context APIs (consumption contract, no AI execution)

`/api/workspaces/{project}/compliance/ai-context/frameworks/{frameworkKey}/releases/{releaseCode}/...` → `summary`, `domains/{domainCode}`, `controls/{controlCode}`, `search`. Returns deterministic AI-ready context; **makes no model call**.

8. Graph APIs

`/api/workspaces/{project}/compliance/graph/frameworks/{frameworkKey}/releases/{releaseCode}` and `.../domains/{domainCode}`, `.../controls/{controlCode}`, `.../requirements/{requirementCode}`.

9. Mapping APIs

`/api/workspaces/{project}/compliance/mappings/...` → `objectives`, `objectives/{objectiveCode}`, `controls/{controlCode}`, `frameworks/compare`, `frameworks/{frameworkKey}/coverage`.

10. Evidence APIs

`/api/workspaces/{project}/compliance/evidence/...` → `types`, `statuses`, `POST context`.

11. Gap APIs

`/api/workspaces/{project}/compliance/gap/...` → `summary`, `domains`, `controls/{controlCode}`, `requirements/{requirementCode}`, `POST context`.

12. Recommendation APIs

/api/workspaces/{project}/compliance/recommendations/... → summary , controls/{controlCode} , requirements/{requirementCode} , POST generate , POST context .

13. Copilot APIs — throttle compliance-copilot

- POST /api/workspaces/{project}/compliance/copilot/message
- GET /api/workspaces/{project}/compliance/copilot/conversations
- GET /api/workspaces/{project}/compliance/copilot/conversations/{conversationUuid}
- POST /api/workspaces/{project}/compliance/copilot/conversations/{conversationUuid}/messages

Example (mock mode):

```
POST /api/workspaces/{project}/compliance/copilot/message
{ "message": "Why is requirement 1-1-1 non-compliant?" }
```

```
{
  "answer": "...cited explanation...",
  "citations": [ { "type": "control", "code": "1-1-1", "source_document_id": "<uuid>" } ],
  "mode": "mock"
}
```

(With AI_COPILOT_DEMO_MODE=true , a demo block with reasoning trace/citations is added. With AI_COPILOT_RAG_ENABLED=true + RAG_ENABLED=true , a rag_context block may be added.)

14. Retrieval / RAG APIs

- **Retrieval (deterministic):** POST /api/workspaces/{project}/compliance/retrieval/query
- **RAG (feature-flagged, returns RAG context only):** POST /api/workspaces/{project}/compliance/rag/query — throttle compliance-rag .

```
POST /api/workspaces/{project}/compliance/rag/query
{ "query": "access control policy" }
```

Returns a bounded, **cited** context package. With RAG disabled, falls back to deterministic retrieval context.

15. Executive APIs (read-only) — throttle compliance-executive

/api/workspaces/{project}/compliance/executive/... → dashboard , scorecard , timeline , explainability , platform . All derived from **real engine data**; no fabricated metrics.

16. Platform AI capability endpoint

GET `/api/ai/platform/capabilities` → modules (adapters), skills, providers, reasoning/retrieval/RAG status, supported contexts, and HUB-wide `module_catalog` (production/reserved/planned).

17. Quenyx AI (Unified AI Workspace — Sprint 20)

v1.0.0 naming: this surface is presented in the UI as **Quenyx AI** (the internal/codename remains *Unified AI Workspace*). API paths are unchanged: routes stay under `/api/ai/*` and the SPA routes stay under `/ai-workspace/*` (the branded `/quenyx-ai/*` path redirects to them). No API contract changed in the rename.

Platform-level AI surface beside Dashboard / Workspaces / Integrations. All routes are flat under `/api/ai/*`, Sanctum-protected, throttled by `ai-workspace`, and **scoped by a required workspace UUID** (query string for reads/deletes, request body for writes) — never a numeric id. Every resource id returned is a UUID. Authorization: `ProjectPolicy::accessAi` (any member) for reads, `administerAi` (owner/admin) for provider/permission changes, plus a fine-grained capability check (`can_use_ai`, `can_manage_templates`, `can_manage_providers`, `can_view_costs`, `can_administer`).

Method	Endpoint	Capability	Notes
GET	/api/ai/workspace/summary	access	counts, tokens, flags + caller permissions
GET	/api/ai/conversations	access	list (metadata only)
POST	/api/ai/conversations	use_ai	start a conversation
GET	/api/ai/conversations/{uuid}	access	conversation + messages
POST	/api/ai/conversations/{uuid}/messages	use_ai	runs the shared AI runtime (mock until AI enabled)
GET	/api/ai/activity	access	AI audit timeline
GET	/api/ai/notifications	access	governance events
GET	/api/ai/usage	view_costs	token usage totals/by-provider/daily
GET	/api/ai/costs	view_costs	tokens × configured pricing; pricing_configured flag
GET	/api/ai/skills	access	dynamic skill catalog (Sprint 19)
GET	/api/ai/capabilities	access	full platform capability catalog
GET	/api/ai/prompt-templates	access	list
POST	/api/ai/prompt-templates	manage_templates	create
PUT	/api/ai/prompt-templates/{uuid}	manage_templates	update
DELETE	/api/ai/prompt-templates/{uuid}	manage_templates	delete
GET	/api/ai/providers	access	provider catalog + prefs (secrets never returned)
PUT	/api/ai/providers/{uuid}/settings	manage_providers + admin	encrypted secret write-only; supports clear_secrets
POST	/api/ai/providers/{uuid}/test	manage_providers + admin	real readiness probe (audited); executable providers run health() , others return status: not_executable
GET	/api/ai/permissions	admin	effective per-role matrix
PUT	/api/ai/permissions	admin	upsert per-role overrides

Cost tracking is **derived** from real `ai_conversations` token counts and an optional `config('ai.workspace.pricing')` table; with no pricing configured, responses carry token totals and `pricing_configured=false` with **no fabricated currency**. Conversation message content is stored only when `ai.feature_flags.prompt_logging` is enabled (privacy-preserving default). Provider settings addressing uses a deterministic UUIDv5 (workspace + provider key) so providers are UUID-addressable even before a settings row exists. Provider id `uuid` is also added (additively) to the workspace list (`GET /api/workspaces`) and `ProjectResource` so the UI can pass the workspace UUID.

Provider catalog & execution (v1.0.0). `GET /api/ai/providers` returns the declarative provider **catalog** (`App\Services\Ai\AiProviderCatalog`) merged with each workspace's saved preferences. Each entry exposes `label`, `type` (`hosted` / `gateway` / `self_hosted` / `custom` / `dev`), declared `capabilities`, `endpoint`, `docs_url`, `is_default`, `executable`, `platform_configured`, `enabled`, `model`, `secret_configured`, `configured`, and `updated_at`. The catalog covers OpenAI, Anthropic Claude, Google Gemini, Azure OpenAI, OpenRouter, Mistral AI, Cohere, xAI Grok, Ollama, LM Studio, vLLM, LiteLLM Gateway, Hugging Face Inference, and a Custom OpenAI-compatible API. A catalogued provider is **configurable but not executable** until a real adapter exists — today only **OpenAI** has a live execution adapter (`executable: true`); all other catalog entries report `executable: false` and **never fabricate connectivity**. The dev-only `mock` provider is **excluded outside local / testing** and is never the production default. The default provider resolves via `AiProviderRegistry::defaultKey()`: an explicit `AI_PROVIDER` wins, otherwise OpenAI when its key is configured, otherwise `mock` only in local/testing, otherwise `''` (an honest *no provider configured* state, surfaced as `summary.has_provider=false`).

18. QynSight Operations Intelligence (Sprint 21)

Workspace-scoped, **UUID-only** AI surface that turns QynSight monitoring data into explainable operational intelligence. All routes are flat under `/api/qynsight/intelligence/*`, Sanctum-protected, throttled by `ai-workspace`, and **scoped by a required workspace UUID** (query string for reads, request body for writes) — never a numeric id. Every entity id in the path/response is a UUID (a deterministic UUIDv5 derived from internal monitoring ids; no schema change, no numeric ids exposed). These endpoints **reuse the Sprint 20 Quenyx AI runtime** (provider registry, prompt orchestration, conversation service, audit) — **no AI logic is duplicated**.

Authorization (every endpoint): `qynsight` module entitlement + monitoring RBAC (`ProjectPolicy::accessAi`) + the per-workspace AI capability `can_use_ai`. Every AI call is audited, provider-logged, and conversation-logged.

Method	Endpoint	Purpose
GET	<code>/api/qynsight/intelligence/overview?workspace={uuid}</code>	Operations Intelligence dashboard: infrastructure health, open alerts, critical services, top operational risks, predicted capacity risks, recent recommendations, recent AI investigations
POST	<code>/api/qynsight/intelligence/copilot</code>	Monitoring Copilot — grounded Q&A; reuses a Quenyx AI conversation (<code>conversation</code> UUID optional to continue a thread)
GET	<code>/api/qynsight/intelligence/recommendations?workspace={uuid}</code>	Evidence-based operational recommendations
POST	<code>/api/qynsight/intelligence/alerts/{uuid}/explain</code>	Alert explanation: impact, most likely causes, evidence used, related alerts, suggested actions, confidence (evidence-derived only)
POST	<code>/api/qynsight/intelligence/alerts/{uuid}/investigate</code>	Deeper alert investigation (explanation + root cause + timeline)
GET	<code>/api/qynsight/intelligence/incidents/{uuid}/timeline?workspace={uuid}</code>	Deterministic incident timeline from actual event timestamps
POST	<code>/api/qynsight/intelligence/hosts/{uuid}/explain</code>	Host (🔥 Explain) — health, performance, capacity narrative
POST	<code>/api/qynsight/intelligence/services/{uuid}/analyze</code>	Service (🔥 Analyze) — why/what changed/impact/action/related
POST	<code>/api/qynsight/intelligence/capacity/{uuid}/predict</code>	Capacity (🔥 Predict) — trend, forecast, exhaustion, action, risk
POST	<code>/api/qynsight/intelligence/infrastructure/{uuid}/impact</code>	Infrastructure Map (🔥 Impact Analysis) — dependencies, SPOF, blast radius

All write endpoints accept `{ "workspace": "<uuid>", ... }`. The copilot accepts `{ "workspace": "<uuid>", "message": "...", "conversation": "<uuid?>" }` and returns `{ "conversation_uuid", "message_uuid", "answer", "evidence" }`. AI narratives carry an `available` / `ai_enabled` flag; when AI is disabled, a clearly flagged mock narrative is returned while the underlying operational data stays real. **No operational data is fabricated**; when evidence is insufficient the response says so.

19. QynAsset Asset Intelligence + AI Adapter Platform (Sprint 22)

AI Adapter discovery (flat under `/api/ai/*`, Sanctum + `throttle:ai-workspace`, required `workspace` UUID, RBAC `accessAi`, audited, **entitlement-filtered** per workspace):

Method	Endpoint	Description
GET	<code>/api/ai/adapters?workspace={uuid}</code>	Every AI module adapter the workspace is entitled to (metadata, capabilities, entities, skills, providers, actions)
GET	<code>/api/ai/adapters/{module}?workspace={uuid}</code>	One adapter descriptor
GET	<code>/api/ai/adapters/capabilities?workspace={uuid}</code>	Aggregated capabilities across entitled adapters
GET	<code>/api/ai/actions?workspace={uuid}</code>	Aggregated contextual actions across entitled adapters

The AI Workspace builds navigation/actions from these — there is **no hard-coded module list**.

QynAsset Asset Intelligence (flat under `/api/qynasset/intelligence/*`, Sanctum + `throttle:ai-workspace`, `qynasset` entitlement, RBAC `accessAi`, `can_use_ai` for AI actions, UUID-only):

Method	Endpoint	Description
GET	<code>/api/qynasset/intelligence/overview?workspace={uuid}</code>	Asset Intelligence dashboard: inventory summary, discovery, capacity, recommendations, recent AI investigations
POST	<code>/api/qynasset/intelligence/copilot</code>	Asset Copilot — grounded Q&A; reuses a Quenyx AI conversation
GET	<code>/api/qynasset/intelligence/recommendations?workspace={uuid}</code>	Evidence-based asset recommendations
POST	<code>/api/qynasset/intelligence/assets/{uuid}/explain</code>	Asset (🔥 Explain) — discovery, activity, hardware, gaps
POST	<code>/api/qynasset/intelligence/assets/{uuid}/dependencies</code>	Dependency (🔥 Analyze) — depends on / serves, neighbors
POST	<code>/api/qynasset/intelligence/assets/{uuid}/lifecycle</code>	Lifecycle (🔥 Forecast) — replacement priority, business impact (warranty/EOL not collected)
POST	<code>/api/qynasset/intelligence/assets/{uuid}/impact</code>	Relationship (🔥 Impact) — blast radius / SPOF
POST	<code>/api/qynasset/intelligence/licenses/review</code>	License (🔥 Review) — honest "not collected" until an inventory/license integration exists

Asset UUIDs are deterministic UUIDv5 (`AssetEntityId`) over the discovered host; **no numeric ids are exposed**. Inventory, hardware, lifecycle, and license facts are **real or honestly absent** — never fabricated.

19b. QynRun — Automation Platform (Sprint 23)

Base `/api/qynrun` . Workspace-scoped, UUID-only. Reads/dry-run require `accessAi` ; approvals, live runs, rollback, and deletes require `administerAi` .

Method	Endpoint	Notes
GET	<code>/automation/adapters</code>	Registered execution adapters; includes <code>live_execution_enabled</code> .
GET	<code>/automation/actions</code>	Action catalog (schema, <code>destructive</code> , <code>rollback</code>).
GET / POST	<code>/automation/workflows</code>	List / create.
GET / PUT / DELETE	<code>/automation/workflows/{uuid}</code>	Read / update / delete.
POST	<code>/automation/workflows/{uuid}/run</code>	<code>mode=dry_run live</code> .
GET / POST	<code>/automation/runbooks</code>	List / create.
GET / PUT / DELETE	<code>/automation/runbooks/{uuid}</code>	Read / update / delete.
POST	<code>/automation/runbooks/{uuid}/run</code>	Run.
GET / POST	<code>/automation/executions</code>	History / ad-hoc dispatch.
GET	<code>/automation/executions/{uuid}</code>	Detail + steps.
POST	<code>/automation/executions/{uuid}/rollback</code>	Rollback a succeeded, rollback-capable run.
POST	<code>/automation/executions/{uuid}/feedback</code>	Operator feedback (learning).
GET	<code>/automation/approvals</code>	Pending approvals.
POST	<code>/automation/approvals/{uuid}/decide</code>	<code>decision=approve reject</code> .
GET	<code>/automation/learning</code>	Aggregated, auditable outcomes.
GET	<code>/intelligence/overview</code>	Automation dashboard payload.
POST	<code>/intelligence/copilot</code>	<code>{ workspace, message, conversation? }</code> — reuses Quenyx AI conversations.
POST	<code>/intelligence/runbooks/suggest</code>	AI-assisted editable runbook draft (never auto-executed).
POST	<code>/intelligence/executions/{uuid}/explain</code>	Explain an execution.

19c. QynReact — Incident Workspace (Sprint 23)

Base `/api/qynreact`. Workspace-scoped, UUID-only, `accessAi`; AI surfaces require `can_use_ai`.

Method	Endpoint	Notes
GET / POST	<code>/incidents</code>	List / open.
GET	<code>/incidents/{uuid}</code>	Unified workspace (timeline, cross-module, automation, ...).
PUT	<code>/incidents/{uuid}</code>	Update status / resolution / postmortem.
POST	<code>/incidents/{uuid}/timeline</code>	Add timeline entry.
POST	<code>/incidents/{uuid}/copilot</code>	<code>{ workspace, message, conversation? }</code> .
POST	<code>/incidents/{uuid}/recommend</code>	Evidence-based response actions.
POST	<code>/incidents/{uuid}/postmortem</code>	Editable postmortem draft.

19d. QynKnow — Enterprise Knowledge Platform (Sprint 24)

Base `/api/qynknow`. Workspace-scoped, UUID-only. Reads require `accessAi`; AI surfaces require `can_use_ai`. All AI goes through `ModuleAiNarrator`; search/timeline/graph are read-models over real rows.

Method	Endpoint	Notes
GET / POST	<code>/documents</code>	List / create knowledge documents (Internal KB).
GET / PUT / DELETE	<code>/documents/{uuid}</code>	Read / update / delete.
GET	<code>/sources</code>	Knowledge Source Registry (operational + planned providers).
GET	<code>/search</code>	Enterprise Search (keyword + semantic) across real rows.
GET	<code>/timeline</code>	Global Timeline (chronological, cross-module read-model).
GET	<code>/graph</code>	Enterprise Knowledge Graph v2 (typed nodes + edges).
POST	<code>/intelligence/explain</code>	Explain a document/entity (editable).
POST	<code>/intelligence/summarize</code>	Summarize (KB/incident/executive/technical).
POST	<code>/intelligence/related</code>	Find related (evidence-based).
POST	<code>/intelligence/draft</code>	Generate KB / runbook draft (editable, never fabricated).
POST	<code>/intelligence/copilot</code>	Knowledge copilot (grounded).

19e. QynSupport — Service Desk + Ticket Intelligence (Sprint 24)

Base `/api/qynsupport`. Workspace-scoped, UUID-only. `accessAi`; AI requires `can_use_ai`.

Method	Endpoint	Notes
GET / POST	<code>/tickets</code>	List / create tickets.
GET / PUT	<code>/tickets/{uuid}</code>	Read / update.
POST	<code>/tickets/{uuid}/intelligence/analyze</code>	Category/priority/impact/assignee/SLA + related (evidence-based).
POST	<code>/tickets/{uuid}/intelligence/copilot</code>	Ticket copilot (grounded).

19f. QynNotify — Notification Intelligence (Sprint 24)

Base `/api/qynnotify`. Workspace-scoped, UUID-only. `accessAi`; AI requires `can_use_ai`.

Method	Endpoint	Notes
GET / POST	<code>/notifications</code>	List (+ correlation groups) / ingest (dedup, correlate, route).
POST	<code>/notifications/{uuid}/read</code>	Mark read.
POST	<code>/intelligence/digest</code>	On-call digest (real active notifications).
POST	<code>/intelligence/executive</code>	Executive summary.
POST	<code>/intelligence/copilot</code>	Notification copilot (grounded).

19g. Collaboration Platform (Sprint 24)

Base `/api/collaboration`. Workspace-scoped, UUID-only, `accessAi`. Polymorphic over any entity (`entity_type`, `entity_uuid`). Shared platform capability — no per-module entitlement.

Method	Endpoint	Notes
GET	<code>/thread</code>	Comments + participants for <code>entity_type</code> + <code>entity_uuid</code> .
POST	<code>/comments</code>	Add comment (optional <code>mentions[]</code> ; mentions become watchers).
POST	<code>/participants</code>	Add <code>watcher</code> / <code>assignee</code> / <code>owner</code> .
DELETE	<code>/participants</code>	Remove a participant role.

19h. QynVA — Enterprise AI Operator + Intelligence surfaces (Sprint 25)

Base `/api/qynva`. Workspace-scoped, UUID-only. Reads require `accessAi`; the operator and executive AI summary require `can_use_ai`; Platform Health and Event Bus introspection require `administerAi`. AI runs through `ModuleAiNarrator` over the **Enterprise Context Engine**;

dashboards/analytics are read-models over real rows (honest `available:false` when data is thin). QynVA proposes editable plans and **never executes**.

Method	Endpoint	Notes
GET	<code>/operator/capabilities</code>	Discovered modules, capabilities, and cross-module action catalog.
POST	<code>/operator/operate</code>	<code>{ workspace, message, conversation? }</code> — reason over enterprise context; editable plan.
GET	<code>/executive</code>	Executive Intelligence dashboard (evidence-based).
POST	<code>/executive/summary</code>	Executive AI summary over the deterministic dashboard.
GET	<code>/analytics</code>	Enterprise Analytics (<code>?days=30</code>): MTTD/MTTR, trends, effectiveness, adoption, KPIs.
GET	<code>/health</code>	Platform Health snapshot (privileged — <code>administerAi</code>).
GET	<code>/events</code>	Platform Event Bus introspection: vocabulary, subscribers, recent events (privileged).

19i. QynBalance — Enterprise Cost Intelligence (Sprint 25)

Base `/api/qynbalance` . Workspace-scoped, UUID-only. Overview requires `accessAi` ; the cost copilot requires `can_use_ai` . Uses **real** platform data; **no fabricated financial values** — monetary figures appear only where `config/cost.php` rates are configured, otherwise counts + "pricing unavailable".

Method	Endpoint	Notes
GET	<code>/cost/overview</code>	Infrastructure cost, license/asset/automation/capacity/cloud sections, budget forecast, recommendations.
POST	<code>/cost/copilot</code>	<code>{ workspace, message, conversation? }</code> — FinOps copilot over deterministic cost evidence.

20. Notes on examples

Example request/response shapes are representative of the controllers' contracts. For exact current fields, call the endpoint against a seeded workspace, or read the corresponding controller/resource in `app/Http/Controllers/**` . Do not assume undocumented fields.