



INTERNAL

Administrator Guide

Quenyx vOPS HUB — v1.0.0 RC1 · Document v2.0

Document Version	2.0
Software Version	v1.0.0 RC1
Applies To	Quenyx vOPS HUB v1.0.0 RC1
Classification	Internal
Owner	Operations
Status	Released
Last Updated	2026-06-29
Document Type	Administrator guide

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0 RC1; native monitoring administration; Unified AI Workspace administration.

Table of Contents

- 1. Login
- 2. Workspace management
- 3. Project management
- 4. Roles
- 5. Users
- 6. Module subscriptions
- 7. QynSight administration
- 8. QynShield administration
- 9. AI settings
- 10. Audit logs
- 11. Integrations
- 12. Limitations
- Quenyx AI administration (Unified AI Workspace — Sprint 20)

12 — Administrator Guide

Audience: Platform administrators. **Scope:** Admin tasks supported by the current product at v1.0.0 RC1 (including the Unified AI Workspace, Sprint 20).

1. Login

- Sign in at the SPA with your credentials (`POST /api/auth/login`). The seeded admin is `admin@quenyx.test` with the password set via `SEED_ADMIN_PASSWORD` at deploy time. **Rotate it after first login.**
- Manage your own profile/password under `auth/me` .

2. Workspace management

- A **workspace** = a **project** (the tenant boundary). Create/list/update/delete via the Projects UI or `/api/workspaces` .
- The seeder provisions **Production Env** and **Staging Env** by default (adjust in `ProjectSeeder`).

3. Project management

- Each workspace has its own monitoring data, compliance scope, members, subscription, audit log, and integrations. Switch the active workspace from the workspace selector.

4. Roles

- Members are attached to a workspace via **memberships** with a **role**. Manage via the Memberships UI or `/api/workspaces/{project}/memberships` (invite, update role, remove).
- Authorization is enforced by `ProjectPolicy` — only members can access a workspace's data.

5. Users

- Users register/are invited, then accept an invite token (`POST /api/invites/{token}/accept`) to join a workspace.

6. Module subscriptions

- View entitlements: `GET /api/workspaces/{project}/entitlements` and `.../modules/access` .
- Manage subscription: `.../subscription` (GET/PUT).
- **Per-module override:** `PUT /api/workspaces/{project}/modules/{moduleKey}/override` — grants/ revokes a module for a workspace. Overrides are **audited**.
- Today the **sidebar shows QynSight only** (other modules are hidden by a frontend flag); module entitlement data still exists for all modules.

7. QynSight administration

- Configure **monitored targets/hosts** (`observe/targets`), **service definitions**, and **alert rules** (create/update/toggle/delete) and **monitoring profiles**.
- Manage **agents**: create enrollment tokens, list/revoke tokens, view metadata, remove agents.
- Ensure the **scheduler** (cron) and **queue worker** run (see Doc 10) — otherwise checks/port scans won't run.

8. QynShield administration

- Requires the **QynShield entitlement** (`project.qynshield`) on the workspace.
- Corpus is seeded by operators (NCA ECC-2:2024 v1). Admins consume `corpus/graph/mapping/evidence/` `gap/recommendation/executive` APIs; there is no in-UI corpus authoring (official-source-only model).

9. AI settings

- AI is configured via **environment variables**, not a UI toggle today (see Doc 10 §11).
- Defaults are safe: mock provider, AI off, no prompt logging, no conversation persistence, RAG off, no tenant-evidence indexing. Change only deliberately.

10. Audit logs

- `GET /api/workspaces/{project}/audit-logs` — review sensitive actions (module overrides, executive reads, copilot, etc.). Use for access reviews and incident investigation.

11. Integrations

- Workspace integrations and billing integrations via `.../integrations` and `.../billing/integrations` . Configure per-workspace settings; secrets are stored as configuration, not in code.

12. Limitations

- No self-service AI toggle UI yet (env-driven).
- QynShield UI beyond executive/demo is roadmap.
- Hidden modules cannot be enabled in the sidebar without changing the frontend flag (out of scope pre-Sprint 20).

Quenyx AI administration (Unified AI Workspace — Sprint 20)

RC1.1: this surface is now branded **Quenyx AI** in the UI (sidebar, header, breadcrumbs). The canonical SPA route stays `/ai-workspace/*` for backward compatibility; the branded `/quenyx-ai/*` path redirects to it. Do not confuse Quenyx AI (the AI control center) with **Workspaces** (tenant/project management).

Open **Quenyx AI** from the top-level sidebar (beside Integrations), then select a workspace. Tabs are grouped into **Workspace** (Overview, Chat, Conversations, History, Activity), **Intelligence** (Skills, Capabilities, Memory, Prompt Templates), **Operations** (Usage, Costs, Providers), and **Administration** (Permissions, Administration, Notifications).

- **Access:** any workspace member can open Quenyx AI; **owner/admin** can administer.
- **Providers** (Operations → Providers): an enterprise provider catalog rendered as cards. Each card shows the provider label, type (hosted / gateway / self-hosted / custom), declared capabilities, endpoint, default indicator, whether a live execution adapter exists (`executable`), whether platform credentials are present, the per-workspace enabled/disabled and configured/secret state, and last-updated. Actions: **Configure/Edit**, **Enable/Disable**, **Test connection** (a real readiness probe — executable providers run a health check, others report "catalog only"), and **Clear secret**. Set the per-workspace model and (optionally) an encrypted API key. Secrets are write-only — stored encrypted and never shown again; the UI only indicates whether a secret is configured. The **default provider** is set by platform configuration (`AI_PROVIDER` , or OpenAI credentials), not per-workspace. The catalog lists OpenAI, Anthropic, Gemini, Azure OpenAI, OpenRouter, Mistral, Cohere, xAI Grok, Ollama, LM Studio, vLLM, LiteLLM, Hugging Face, and a Custom OpenAI-compatible API; today **only OpenAI is executable** — others are configurable but not yet executable, and the **mock** provider is hidden outside local/testing. AI execution still requires the platform AI flags.
- **Permissions** (Administration → Permissions): per-role matrix (`Use AI` , `Manage templates` , `Manage providers` , `View costs` , `Administer`). Rows are additive overrides on top of role defaults; the `owner` row is always full and locked.
- **Cost tracking:** amounts appear only when pricing is configured in `config/ai.php` (`ai.workspace.pricing`); otherwise token usage is shown without monetary values.
- **Audit:** provider/template/permission changes and conversations are recorded in the audit log and surfaced under Activity / Notifications.
- **Disable:** set `AI_WORKSPACE_ENABLED=false` to hide the surface (returns 404 from the API).