QUENYX VOPS HUB

INTERNAL

Administrator Guide

Quenyx vOPS HUB — v1.0.0 RC1 · Document v2.1

Document Version	2.1
Software Version	v1.0.0 RC1
Applies To	Quenyx vOPS HUB v1.0.0 RC1
Classification	Internal
Owner	Operations
Status	Released
Last Updated	2026-06-30
Document Type	Administrator guide

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0 RC1; native monitoring administration; Unified AI Workspace administration.
2.1	2026-06-30	Added administration of QynSight Operations Intelligence (Sprint 21): entitlement + `can_use_ai` capability, audit, and rate limits.

Table of Contents

12 — Administrator Guide

Audience: Platform administrators. **Scope:** Admin tasks supported by the current product at v1.0.0 RC1 (including the Unified AI Workspace, Sprint 20).

---

1. Login

- Sign in at the SPA with your credentials (`POST /api/auth/login`). The seeded admin is `admin@quenyx.test` with the password set via `SEED_ADMIN_PASSWORD` at deploy time. **Rotate it after first login.**
- Manage your own profile/password under `auth/me` .

2. Workspace management

- A **workspace** = a **project** (the tenant boundary). Create/list/update/delete via the Projects UI or `/api/workspaces` .
- The seeder provisions **Production Env** and **Staging Env** by default (adjust in `ProjectSeeder`).

3. Project management

- Each workspace has its own monitoring data, compliance scope, members, subscription, audit log, and integrations. Switch the active workspace from the workspace selector.

4. Roles

- Members are attached to a workspace via **memberships** with a **role**. Manage via the Memberships UI or `/api/workspaces/{project}/memberships` (invite, update role, remove).
- Authorization is enforced by `ProjectPolicy` — only members can access a workspace's data.

5. Users

- Users register/are invited, then accept an invite token (`POST /api/invites/{token}/accept`) to join a workspace.

6. Module subscriptions

- View entitlements: `GET /api/workspaces/{project}/entitlements` and `…/modules/access` .
- Manage subscription: `…/subscription` (GET/PUT).
- **Per-module override:** `PUT /api/workspaces/{project}/modules/{moduleKey}/override` — grants/ revokes a module for a workspace. Overrides are **audited**.
- Today the **sidebar shows QynSight only** (other modules are hidden by a frontend flag); module entitlement data still exists for all modules.

7. QynSight administration

- Configure **monitored targets/hosts** (`observe/targets`), **service definitions**, and **alert rules** (create/update/toggle/delete) and **monitoring profiles**.
- Manage **agents**: create enrollment tokens, list/revoke tokens, view metadata, remove agents.
- Ensure the **scheduler** (cron) and **queue worker** run (see Doc 10) — otherwise checks/port scans won't run.

8. QynShield administration

- Requires the **QynShield entitlement** (`project.qynshield`) on the workspace.
- Corpus is seeded by operators (NCA ECC-2:2024 v1). Admins consume corpus/graph/mapping/evidence/ gap/recommendation/executive APIs; there is no in-UI corpus authoring (official-source-only model).

9. AI settings

- AI is configured via **environment variables**, not a UI toggle today (see Doc 10 §11).
- Defaults are safe: mock provider, AI off, no prompt logging, no conversation persistence, RAG off, no tenant-evidence indexing. Change only deliberately.

10. Audit logs

- `GET /api/workspaces/{project}/audit-logs` — review sensitive actions (module overrides, executive reads, copilot, etc.). Use for access reviews and incident investigation.

11. Integrations

- Workspace integrations and billing integrations via `…/integrations` and `…/billing/integrations` . Configure per-workspace settings; secrets are stored as configuration, not in code.

12. Limitations

- No self-service AI toggle UI yet (env-driven).
- QynShield UI beyond executive/demo is roadmap.
- Hidden modules cannot be enabled in the sidebar without changing the frontend flag (out of scope pre-Sprint 20).

Quenyx AI administration (Unified AI Workspace — Sprint 20)

> **RC1.1:** this surface is now branded **Quenyx AI** in the UI (sidebar, header, breadcrumbs). The canonical SPA route stays `/ai-workspace/*` for backward compatibility; the branded `/quenyx-ai/*` path redirects to it. Do not confuse Quenyx AI (the AI control center) with **Workspaces** (tenant/project management).

Open **Quenyx AI** from the top-level sidebar (beside Integrations), then select a workspace. Tabs are grouped into **Workspace** (Overview, Chat, Conversations, History, Activity), **Intelligence** (Skills, Capabilities, Memory, Prompt Templates), **Operations** (Usage, Costs, Providers), and **Administration** (Permissions, Administration, Notifications).

- **Access**: any workspace member can open Quenyx AI; **owner/admin** can administer.
- **Providers** (Operations → Providers): an enterprise provider catalog rendered as cards. Each card shows the provider label, type (hosted / gateway / self-hosted / custom), declared capabilities, endpoint, default indicator, whether a live execution adapter exists (`executable`), whether platform credentials are present, the per-workspace enabled/disabled and configured/secret state, and last-updated. Actions: **Configure/Edit**, **Enable/Disable**, **Test connection** (a real readiness probe — executable providers run a health check, others report "catalog only"), and **Clear secret**. Set the per-workspace model and (optionally) an encrypted API key. Secrets are write-only — stored encrypted and never shown again; the UI only indicates whether a secret is configured. The **default provider** is set by platform configuration (`AI_PROVIDER` , or OpenAI credentials), not per-workspace. The catalog lists OpenAI, Anthropic, Gemini, Azure OpenAI, OpenRouter, Mistral, Cohere, xAI Grok, Ollama, LM Studio, vLLM, LiteLLM, Hugging Face, and a Custom OpenAI-compatible API; today **only OpenAI is executable** — others are configurable but not yet executable, and the **mock** provider is hidden outside local/testing. AI execution still requires the platform AI flags.
- **Permissions** (Administration → Permissions): per-role matrix (`Use AI` , `Manage templates` , `Manage providers` , `View costs` , `Administer`). Rows are additive overrides on top of role defaults; the `owner` row is always full and locked.
- **Cost tracking**: amounts appear only when pricing is configured in `config/ai.php` (`ai.workspace.pricing`); otherwise token usage is shown without monetary values.
- **Audit**: provider/template/permission changes and conversations are recorded in the audit log and surfaced under Activity / Notifications.
- **Disable**: set `AI_WORKSPACE_ENABLED=false` to hide the surface (returns 404 from the API).

QynSight Operations Intelligence administration (Sprint 21)

Operations Intelligence adds an explainable AI layer to QynSight (Monitoring Copilot, Alert/Root-Cause/Capacity/Performance/Infrastructure/Service-Health intelligence, evidence-based recommendations, and an Operations Intelligence dashboard). It **reuses the Quenyx AI runtime** — there is nothing new to provision, no new provider to configure, and no new secret to manage.

- **Entitlement & capability**: a workspace must have the `qynsight` module entitlement *and* the `can_use_ai` AI capability (manage it in Quenyx AI → Administration → Permissions per role). A member also needs monitoring RBAC. Without these, the endpoints return `403 Locked`.
- **AI posture**: governed by the **same** AI flags as Quenyx AI (§9 and Doc 10 §11). With AI disabled, the Copilot/✨ actions return a **clearly flagged mock narrative** over **real** monitoring evidence — no external model calls and no fabricated operational data.
- **Where admins see it**: the Operations Intelligence dashboard lives under QynSight (`/observe/operations-intelligence`); contextual ✨ **Quenyx AI** actions appear on the Hosts, Services, Alerts, Capacity Planning, and Infrastructure Map pages.
- **Audit & limits**: every AI action is audited (`audit_logs`, `action LIKE 'ai%'`), provider-logged, conversation-logged, and rate limited via `throttle:ai-workspace`. Copilot threads are real Quenyx AI conversations and appear in the AI Activity/History surfaces.
- **Data prerequisites**: insights are only as rich as the monitoring data — ensure the scheduler and agents are healthy so hosts/services/alerts/metrics/capacity are current.

QynAsset Asset Intelligence + AI Adapter Platform (Sprint 22)

- **Entitlement**: Asset Intelligence requires the `qynasset` module entitlement for the workspace, on top of the same RBAC (`accessAi`) and `can_use_ai` capability used by all AI actions.
- **Where**: the **Asset Intelligence** dashboard lives at `/app/workspaces/:id/qynasset/intelligence`; contextual ✨ actions (Explain / Analyze / Forecast / Impact / Review) open the Asset Copilot, which reuses Quenyx AI conversations.
- **Data prerequisites**: an asset is a **discovered host** — enroll agents and define monitored hosts so inventory, agent status, and hardware facts are present. Software-license and warranty/EOL data are **not collected** until an inventory/license integration is connected; until then those capabilities honestly report "not collected" rather than fabricating values.
- **Adapter discovery**: administrators/integrators can inspect which AI modules a workspace can use via `GET /api/ai/adapters` (entitlement-filtered). Every AI action remains audited, provider-logged, conversation-logged, and rate-limited.

Automation Platform administration (Sprint 23)

QynRun/QynReact run on the shared Automation Platform. Administration is **safe by default**:

- **Live execution master switch.** `automation.live_execution` (env `AUTOMATION_LIVE_EXECUTION`) defaults to **OFF**. While OFF, every action is dry-run only — no side effects are possible.
- **Per-runner enable flags.** Script, SSH, and PowerShell runners each have their own enable flag in `config/automation.php`; until enabled they honestly report `skipped` for live requests.
- **HTTP allowlist.** REST/Webhook actions only reach hosts in `automation.allowed_hosts`.
- **Approval gate.** No destructive/live action runs without human approval. Approving, rejecting, rolling back, and deleting require `administerAi`; viewing and dry-run require `accessAi`.
- **Rollback.** Successful, rollback-capable executions can be undone from the Executions view.

- **Auditing & learning.** All automation events are written via `AutomationAuditLogger`; outcomes are captured as auditable learning records (no model training, no hidden state) and inform — but never auto-trigger — AI recommendations.

All automation and incident data is workspace-isolated and UUID-only. See the **Automation Platform Guide (Doc 24)** and **Incident Response Guide (Doc 27)**.