



QUENYX

V O P S H U B

INTERNAL

Administrator Guide

Quenyx vOPS HUB — v1.0.0 RC1 · Document v2.1

Document Version	2.1
Software Version	v1.0.0 RC1
Applies To	Quenyx vOPS HUB v1.0.0 RC1
Classification	Internal
Owner	Operations
Status	Released
Last Updated	2026-06-30
Document Type	Administrator guide

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0 RC1; native monitoring administration; Unified AI Workspace administration.
2.1	2026-06-30	Added administration of QynSight Operations Intelligence (Sprint 21): entitlement + `can_use_ai` capability, audit, and rate limits.

Table of Contents

1. Login

2. Workspace management

3. Project management

4. Roles

5. Users

6. Module subscriptions

7. QynSight administration

8. QynShield administration

9. AI settings

10. Audit logs

11. Integrations

12. Limitations

Quenyx AI administration (Unified AI Workspace — Sprint 20)

QynSight Operations Intelligence administration (Sprint 21)

QynAsset Asset Intelligence + AI Adapter Platform (Sprint 22)

12 — Administrator Guide

Audience: Platform administrators. **Scope:** Admin tasks supported by the current product at v1.0.0 RC1 (including the Unified AI Workspace, Sprint 20).

1. Login

- Sign in at the SPA with your credentials (`POST /api/auth/login`). The seeded admin is `admin@quenyx.test` with the password set via `SEED_ADMIN_PASSWORD` at deploy time. **Rotate it after first login.**
- Manage your own profile/password under `auth/me` .

2. Workspace management

- A **workspace** = a **project** (the tenant boundary). Create/list/update/delete via the Projects UI or `/api/workspaces` .
- The seeder provisions **Production Env** and **Staging Env** by default (adjust in `ProjectSeeder`).

3. Project management

- Each workspace has its own monitoring data, compliance scope, members, subscription, audit log, and integrations. Switch the active workspace from the workspace selector.

4. Roles

- Members are attached to a workspace via **memberships** with a **role**. Manage via the Memberships UI or `/api/workspaces/{project}/memberships` (invite, update role, remove).
- Authorization is enforced by `ProjectPolicy` — only members can access a workspace's data.

5. Users

- Users register/are invited, then accept an invite token (`POST /api/invites/{token}/accept`) to join a workspace.

6. Module subscriptions

- View entitlements: `GET /api/workspaces/{project}/entitlements` and `.../modules/access` .
- Manage subscription: `.../subscription` (GET/PUT).
- **Per-module override:** `PUT /api/workspaces/{project}/modules/{moduleKey}/override` — grants/ revokes a module for a workspace. Overrides are **audited**.
- Today the **sidebar shows QynSight only** (other modules are hidden by a frontend flag); module entitlement data still exists for all modules.

7. QynSight administration

- Configure **monitored targets/hosts** (`observe/targets`), **service definitions**, and **alert rules** (create/update/toggle/delete) and **monitoring profiles**.
- Manage **agents**: create enrollment tokens, list/revoke tokens, view metadata, remove agents.
- Ensure the **scheduler** (cron) and **queue worker** run (see Doc 10) — otherwise checks/port scans won't run.

8. QynShield administration

- Requires the **QynShield entitlement** (`project.qynshield`) on the workspace.
- Corpus is seeded by operators (NCA ECC-2:2024 v1). Admins consume `corpus/graph/mapping/evidence/` `gap/recommendation/executive` APIs; there is no in-UI corpus authoring (official-source-only model).

9. AI settings

- AI is configured via **environment variables**, not a UI toggle today (see Doc 10 §11).
- Defaults are safe: mock provider, AI off, no prompt logging, no conversation persistence, RAG off, no tenant-evidence indexing. Change only deliberately.

10. Audit logs

- `GET /api/workspaces/{project}/audit-logs` — review sensitive actions (module overrides, executive reads, copilot, etc.). Use for access reviews and incident investigation.

11. Integrations

- Workspace integrations and billing integrations via `.../integrations` and `.../billing/integrations` . Configure per-workspace settings; secrets are stored as configuration, not in code.

12. Limitations

- No self-service AI toggle UI yet (env-driven).
- QynShield UI beyond executive/demo is roadmap.
- Hidden modules cannot be enabled in the sidebar without changing the frontend flag (out of scope pre-Sprint 20).

Quenyx AI administration (Unified AI Workspace — Sprint 20)

RC1.1: this surface is now branded **Quenyx AI** in the UI (sidebar, header, breadcrumbs). The canonical SPA route stays `/ai-workspace/*` for backward compatibility; the branded `/quenyx-ai/*` path redirects to it. Do not confuse Quenyx AI (the AI control center) with **Workspaces** (tenant/project management).

Open **Quenyx AI** from the top-level sidebar (beside Integrations), then select a workspace. Tabs are grouped into **Workspace** (Overview, Chat, Conversations, History, Activity), **Intelligence** (Skills, Capabilities, Memory, Prompt Templates), **Operations** (Usage, Costs, Providers), and **Administration** (Permissions, Administration, Notifications).

- **Access:** any workspace member can open Quenyx AI; **owner/admin** can administer.
- **Providers** (Operations → Providers): an enterprise provider catalog rendered as cards. Each card shows the provider label, type (hosted / gateway / self-hosted / custom), declared capabilities, endpoint, default indicator, whether a live execution adapter exists (`executable`), whether platform credentials are present, the per-workspace enabled/disabled and configured/secret state, and last-updated. Actions: **Configure/Edit**, **Enable/Disable**, **Test connection** (a real readiness probe — executable providers run a health check, others report "catalog only"), and **Clear secret**. Set the per-workspace model and (optionally) an encrypted API key. Secrets are write-only — stored encrypted and never shown again; the UI only indicates whether a secret is configured. The **default provider** is set by platform configuration (`AI_PROVIDER` , or OpenAI credentials), not per-workspace. The catalog lists OpenAI, Anthropic, Gemini, Azure OpenAI, OpenRouter, Mistral, Cohere, xAI Grok, Ollama, LM Studio, vLLM, LiteLLM, Hugging Face, and a Custom OpenAI-compatible API; today **only OpenAI is executable** — others are configurable but not yet executable, and the **mock** provider is hidden outside local/testing. AI execution still requires the platform AI flags.
- **Permissions** (Administration → Permissions): per-role matrix (`Use AI` , `Manage templates` , `Manage providers` , `View costs` , `Administer`). Rows are additive overrides on top of role defaults; the `owner` row is always full and locked.
- **Cost tracking:** amounts appear only when pricing is configured in `config/ai.php` (`ai.workspace.pricing`); otherwise token usage is shown without monetary values.
- **Audit:** provider/template/permission changes and conversations are recorded in the audit log and surfaced under Activity / Notifications.
- **Disable:** set `AI_WORKSPACE_ENABLED=false` to hide the surface (returns 404 from the API).

QynSight Operations Intelligence administration (Sprint 21)

Operations Intelligence adds an explainable AI layer to QynSight (Monitoring Copilot, Alert/Root-Cause/Capacity/Performance/Infrastructure/Service-Health intelligence, evidence-based recommendations, and an Operations Intelligence dashboard). It **reuses the Quenyx AI runtime** — there is nothing new to provision, no new provider to configure, and no new secret to manage.

- **Entitlement & capability:** a workspace must have the `qynsight` module entitlement *and* the `can_use_ai` AI capability (manage it in Quenyx AI → Administration → Permissions per role). A member also needs monitoring RBAC. Without these, the endpoints return `403 Locked`.
- **AI posture:** governed by the **same** AI flags as Quenyx AI (§9 and Doc 10 §11). With AI disabled, the Copilot/ ✨ actions return a **clearly flagged mock narrative** over **real** monitoring evidence — no external model calls and no fabricated operational data.
- **Where admins see it:** the Operations Intelligence dashboard lives under QynSight (`/observe/operations-intelligence`); contextual ✨ **Quenyx AI** actions appear on the Hosts, Services, Alerts, Capacity Planning, and Infrastructure Map pages.
- **Audit & limits:** every AI action is audited (`audit_logs` , `action LIKE 'ai%'`), provider-logged, conversation-logged, and rate limited via `throttle:ai-workspace` . Copilot threads are real Quenyx AI conversations and appear in the AI Activity/History surfaces.
- **Data prerequisites:** insights are only as rich as the monitoring data — ensure the scheduler and agents are healthy so hosts/services/alerts/metrics/capacity are current.

QynAsset Asset Intelligence + AI Adapter Platform (Sprint 22)

- **Entitlement:** Asset Intelligence requires the `qynasset` module entitlement for the workspace, on top of the same RBAC (`accessAi`) and `can_use_ai` capability used by all AI actions.
- **Where:** the **Asset Intelligence** dashboard lives at `/app/workspaces/:id/qynasset/intelligence` ; contextual ✨ actions (Explain / Analyze / Forecast / Impact / Review) open the Asset Copilot, which reuses Quenyx AI conversations.
- **Data prerequisites:** an asset is a **discovered host** — enroll agents and define monitored hosts so inventory, agent status, and hardware facts are present. Software-license and warranty/EOL data are **not collected** until an inventory/license integration is connected; until then those capabilities honestly report "not collected" rather than fabricating values.
- **Adapter discovery:** administrators/integrators can inspect which AI modules a workspace can use via `GET /api/ai/adapters` (entitlement-filtered). Every AI action remains audited, provider-logged, conversation-logged, and rate-limited.