



PUBLIC / EXTERNAL

AI User Guide

Quenyx vOPS HUB — v1.0.0 RC1 · Document v2.0

Document Version	2.0
Software Version	v1.0.0 RC1
Applies To	Quenyx vOPS HUB v1.0.0 RC1
Classification	Public / External
Owner	Product
Status	Released
Last Updated	2026-06-29
Document Type	User guide

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0 RC1; Quenyx AI as a shared platform; Unified AI Workspace.

Table of Contents

- | 1. What Quenyx AI is
- | 2. What it can do today ● / ●
- | 3. What is feature-flagged ●
- | 4. How citations work
- | 5. How reasoning works
- | 6. What AI cannot do
- | 7. Mock mode vs AI mode
- | 8. Prompt logging policy
- | 9. Privacy / security
- | 10. Safe usage guidance
- | 11. Demo prompts

16 — AI User Guide

Audience: All users. **Principle:** Quenyx AI is **governable, deterministic-first, cited**, and **off by default**.

1. What Quenyx AI is

A **shared AI platform** behind the HUB that explains compliance and operational data — not a free-form chatbot. Business logic (deterministic engines) decides *what* is true; AI only helps *phrase* it, always with citations.

RC1.1: the workspace AI control center (Sprint 20's *Unified AI Workspace*) is now branded **Quenyx AI** and opens from the top-level sidebar (beside Integrations). It is distinct from **Workspaces** (tenant/project management). Tabs are grouped into Workspace, Intelligence, Operations, and Administration. Under **Operations** → **Providers** you can browse the provider catalog (OpenAI, Anthropic, Gemini, Azure OpenAI, OpenRouter, Mistral, Cohere, xAI Grok, Ollama, LM Studio, vLLM, LiteLLM, Hugging Face, Custom), configure credentials (write-only/encrypted), enable/disable, and run a real **Test connection**. Providers without a live adapter are clearly marked "catalog only"; only **OpenAI** executes today.

2. What it can do today /

- **Compliance Copilot** — answers questions about the loaded framework (NCA ECC-2:2024) with **citations**, backed by deterministic reasoning. (Mock mode by default; real-model when enabled.)
- **Skills** — corpus search, knowledge graph, framework mapping, evidence, gap assessment, recommendations — all reusing deterministic services.
- **Capability catalog** — `GET /api/ai/platform/capabilities` shows exactly what the platform can do right now.

3. What is feature-flagged

- **Real-model AI** (`AI_ENABLED` / `AI_PROVIDER=openai`) — off by default.
- **RAG** (`RAG_ENABLED` , `AI_COPILOT_RAG_ENABLED` , `EMBEDDINGS_ENABLED`) — off; metadata-only with deterministic fallback when on.
- **Demo mode** (`AI_COPILOT_DEMO_MODE`) — surfaces reasoning trace + citations + sources.

4. How citations work

Every compliance answer references the **official source** (control/requirement + source document). If the platform cannot cite a source, it **does not answer**. This is the anti-hallucination guarantee.

5. How reasoning works

A deterministic **Reasoning Engine** applies explicit rules to the corpus + evidence and produces a `ReasoningOutput` (findings, fired rule IDs, explanation). The same inputs always yield the same result. AI renders this output; it does not invent it.

6. What AI cannot do

- It **cannot** answer without a citable source.
- It **cannot** access tenant evidence embeddings (not indexed by default).
- It **cannot** call a model unless an operator enabled it (`AI_ENABLED`).
- It is **not** an autonomous agent and does **not** take actions on your infrastructure.

7. Mock mode vs AI mode

	Mock mode (default)	AI mode (<code>AI_ENABLED=true</code>)
Model calls	none	OpenAI provider
Output	deterministic mock phrasing	model-phrased, still cited + reasoning-gated
Safety	maximal (no external calls)	governed by the same guardrails

Mock mode is safe for demos and produces real, cited structure without contacting any model.

8. Prompt logging policy

Off by default. Prompt/response **content is not stored** unless an operator sets `AI_PROMPT_LOGGING_ENABLED=true` .

9. Privacy / security

- Conversation persistence **off by default**.
- Tenant **evidence is never embedded** by default.
- All AI routes are authenticated, entitlement-gated, and rate-limited.
- Models/keys come from server env, never hardcoded.

10. Safe usage guidance

- Scope questions to a loaded framework (NCA ECC-2:2024) for best results.
- Treat answers as **explanations of platform data**, verifiable via their citations.
- If an answer lacks a citation, it won't be given — refine the question or check entitlement/scope.

11. Demo prompts

- "Summarize domain 1 of NCA ECC-2:2024."
- "Why is requirement 1-1-1 non-compliant?"
- "What evidence is missing for control 2-1-1?"
- "What do you recommend for the open gaps in domain 2?"
- "Show the capability catalog for the AI platform." (→ `GET /api/ai/platform/capabilities`)