

INTERNAL — OPERATIONS

Operations Runbook

Quenyx vOPS HUB — v1.0.0 · Document v2.1

Document Version	2.1
Software Version	v1.0.0
Applies To	Quenyx vOPS HUB v1.0.0
Classification	Internal — Operations
Owner	Operations / SRE
Status	Released
Last Updated	2026-06-30
Document Type	Operations runbook

REVISION HISTORY

Version	Date	Notes
1.0	2026	Initial v1 pack (through Sprint 19).
2.0	2026-06-29	Aligned to v1.0.0; native monitoring operations; Unified AI Workspace operations. See also <code>`docs/OBSERVE_RUNBOOK.md`</code> .
2.1	2026-06-30	Added Operations Intelligence (Sprint 21) operations — shares the Quenyx AI runtime; no new services to operate.

Table of Contents

- 1. Daily checks
- 2. Weekly checks
- 3. Backups
- 4. Restores
- 5. Logs
- 6. Queues
- 7. Scheduler
- 8. SSL
- 9. DB health
- 10. Cache
- 11. Incident handling
- 12. Rollback
- 13. Emergency: disable AI
- 14. Emergency: disable a module
- 15. Troubleshooting commands
- Quenyx AI (Unified AI Workspace — Sprint 20) — operations
- QynSight Operations Intelligence (Sprint 21) — operations
- QynAsset Asset Intelligence + AI Adapter Platform (Sprint 22)
- Automation & Incident operations (Sprint 23)
- Knowledge, Service Desk & Notification operations (Sprint 24)
- Enterprise Intelligence operations (Sprint 25)

18 — Operations Runbook

Audience: Ops / SRE. **Scope:** Run, monitor, and recover a Quenyx vOPS HUB deployment. Commands assume the backend dir and a Linux host (adjust paths).

1. Daily checks

- ☐ `curl -s https://<host>/api/health` → healthy; gateway `/health` → `{"status":"ok"}`.
- ☐ Scheduler heartbeat: `tail backend/storage/logs/scheduler.log` shows recent `observe:run-checks`.
- ☐ Queue worker alive: `systemctl status quenyx-queue`.
- ☐ Error log scan: `tail -n 200 backend/storage/logs/laravel.log` for ERROR/exception spikes.
- ☐ Disk and DB connectivity OK.

2. Weekly checks

- ☐ Review `audit-logs` for unexpected privileged actions.
- ☐ Verify backups exist and are restorable (spot-restore to staging).
- ☐ Check SSL certificate expiry.
- ☐ Review failed jobs: `php artisan queue:failed`.
- ☐ Confirm AI flags still match intended posture (see §11/§12 below).

3. Backups

- **DB:** `mysqldump -u <user> -p <db> > backup-$(date +%F).sql` (automate nightly, encrypt, offsite).
- **Storage:** archive `backend/storage/` (agent binaries, artifacts).
- **Secrets:** `.env` lives in a secrets manager, not in backups-in-the-clear.

4. Restores

- **DB:** `mysql -u <user> -p <db> < backup-YYYY-MM-DD.sql`.
- **Storage:** extract the archive to `backend/storage/`.
- Re-run `php artisan config:cache` after restore; verify health.

5. Logs

- Laravel: `backend/storage/logs/laravel.log`; scheduler: `scheduler.log`.
- Gateway: service logs via `journalctl -u quenyx-gateway`.
- Nginx: access/error logs under `/var/log/nginx/`.

6. Queues

- Status: `systemctl status quenyx-queue` ; restart: `systemctl restart quenyx-queue` .
- Retry failed: `php artisan queue:retry all` ; inspect: `php artisan queue:failed` .
- Required for **port scans** and **RAG indexing jobs**.

7. Scheduler

- Cron must run `php artisan schedule:run` every minute (`crontab -u www-data -l`).
- If monitoring shows "never": confirm cron, PHP path, and `storage/logs` writability.

8. SSL

- Renew certs (e.g. certbot); reload Nginx: `systemctl reload nginx` .
- Verify HTTP→HTTPS redirect and HSTS as configured.

9. DB health

- `php artisan migrate:status` → all applied.
- Monitor connections, slow queries, and disk. Alert on connectivity loss.

10. Cache

- Clear after config changes: `php artisan config:clear` (dev) / `config:cache` (prod).
- Route/view caches: `php artisan route:cache` , `view:cache` . Clear all: `php artisan optimize:clear` .

11. Incident handling

1. Confirm scope via health endpoints + logs.
2. Identify the failing component (backend/gateway/db/queue/scheduler).
3. Mitigate (restart service, fail over, disable a feature flag).
4. Capture logs and audit entries.
5. Post-incident: file the fix forward (new migration/code), update this runbook.

12. Rollback

- **App**: deploy previous tag; rebuild backend/frontend/gateway; restart.
- **DB**: restore from pre-deploy dump (don't `migrate:rollback` destructive changes in prod).
- **Corpus**: roll back to the previous active revision (deterministic snapshot).

13. Emergency: disable AI

```
# In backend/.env
AI_ENABLED=false
AI_PROVIDER=mock
RAG_ENABLED=false
AI_COPILOT_RAG_ENABLED=false
php artisan config:cache # or config:clear in dev
```

This forces mock mode — **no external model calls** — while keeping deterministic features working.

14. Emergency: disable a module

- Per workspace: revoke via `PUT /api/workspaces/{project}/modules/{moduleKey}/override` (audited).
- QynSight gate: remove the `qynsight` entitlement to disable `observe/*` for a workspace.
- Frontend: a hidden module stays hidden via the existing flag (no action needed).

15. Troubleshooting commands

```
php artisan about # environment summary (needs mbstring)
php artisan route:list # verify routes/no collisions
php artisan migrate:status # migration state
php artisan queue:failed # failed jobs
php artisan optimize:clear # clear caches
tail -f backend/storage/logs/laravel.log
journalctl -u quenyx-gateway -f
```

Quenyx AI (Unified AI Workspace — Sprint 20) — operations

v1.0.0: UI label is **Quenyx AI**. Routes unchanged (`/api/ai/*` , `/ai-workspace/*` ; `/quenyx-ai/*` redirects). The provider list is now catalog-driven and the dev-only **mock** provider is hidden outside `local / testing` ; the production default provider is **OpenAI** when configured, otherwise an honest "no provider configured" state (never mock).

Deploy (after pulling):

```
cd backend
composer install --no-dev --optimize-autoloader # or: composer dump-autoloader -o
php artisan migrate --force # adds projects.uuid (backfilled),
                             # ai_prompt_templates, ai_provider_settings,
php artisan route:list | grep ai # expect /api/ai/workspace/summary, conversat
php artisan config:cache
```

- **Enable/disable** the surface: `AI_WORKSPACE_ENABLED` (default `true`). When `false`, all `/api/ai/*` workspace endpoints return 404 and the sidebar item leads to an empty surface.
- **Safe by default:** with `AI_ENABLED=false` (default), chat uses the mock provider; nothing reaches an external model. No raw provider secrets are stored — `ai_provider_settings.settings` is encrypted (depends on a valid `APP_KEY`; rotating `APP_KEY` invalidates stored secrets, which must be re-entered).
- **Default provider (v1.0.0):** leave `AI_PROVIDER` unset to let the registry resolve it safely (OpenAI when `OPENAI_API_KEY` is set; `mock` only in local/testing; otherwise none). Set `AI_PROVIDER` explicitly to override. The provider **catalog** is informational; only providers with a real adapter (today OpenAI) are executable. The **Test connection** action / `POST /api/ai/providers/{uuid}/test` runs a genuine `health()` for executable providers and is audited.
- **Cost tracking** shows currency only when `ai.workspace.pricing` is configured; otherwise it is token-only by design (no fabricated amounts).
- **Auditing:** AI conversations, provider/template/permission changes are written to `audit_logs` (`action LIKE 'ai%'`); surface them via the Activity/Notifications tabs or query the table directly.
- **Rate limiting:** `throttle:ai-workspace` (default 120/min, `AI_WORKSPACE_RATE_LIMIT`).

QynSight Operations Intelligence (Sprint 21) — operations

Operations Intelligence is an **additive layer** that reuses the Quenyx AI runtime — there is **no new daemon, queue, or service to operate**. It ships with the same backend deploy (no new migrations; UUIDs are derived deterministically at runtime). After pulling:

```
cd backend
composer dump-autoload -o
php artisan route:list | grep qynsight/intelligence # expect overview, copilot, alerts/.../exp
php artisan config:cache
```

- **Prerequisites:** the workspace must have the `qynsight` entitlement and the `can_use_ai` AI capability; the requesting member needs monitoring RBAC. Live operational data still requires the **scheduler** (`observe:run-checks`, `observe:evaluate-alerts`) and the **queue worker** (port scans) per §6/§7 — Operations Intelligence reads that data and never fabricates it.
- **AI posture:** governed by the same flags as Quenyx AI. With `AI_ENABLED=false` the Copilot and ✨ actions return a **clearly flagged mock narrative** over **real** evidence; the emergency "disable AI" steps in §13 also apply here (no external model calls).
- **Endpoints:** flat under `/api/qynsight/intelligence/*`, Sanctum + `throttle:ai-workspace` (shares the same limiter/budget as Quenyx AI).
- **Auditing:** alert investigations and other AI actions write to `audit_logs` (`action LIKE 'ai%'`); recent investigations also surface on the Operations Intelligence dashboard.
- **Troubleshooting:** a `403 Locked` means the `qynsight` entitlement or `can_use_ai` capability is missing; an empty/limited narrative with "insufficient evidence" means monitoring has not yet collected enough data (check scheduler/agents), **not** an AI failure.

QynAsset Asset Intelligence + AI Adapter Platform (Sprint 22)

- **Endpoints:** flat under `/api/qynasset/intelligence/*`, Sanctum + `throttle:ai-workspace` (shares the Quenyx AI limiter/budget); adapter discovery under `/api/ai/adapters` + `/api/ai/actions`.
 - **Entitlement/RBAC:** a `403` means the `qynasset` entitlement, `accessAi` RBAC, or `can_use_ai` capability is missing. UUID-only — a `404 Asset not found` means the asset UUID does not resolve to a host in this workspace.
 - **Empty/honest output:** an asset overview with `total: 0` means no hosts/agents are discovered yet (enroll agents / define hosts) — not an AI failure. License/lifecycle-date sections returning `available: false` is **expected** (no inventory/license integration configured), **not** an error; Quenyx AI never fabricates those facts.
 - **Auditing:** asset AI actions write to `audit_logs` (`action` LIKE `'asset_intelligence_%'` and `ai_adapter_discovery_%`); recent investigations also surface on the Asset Intelligence dashboard.
 - **No duplicated AI:** QynAsset reuses the shared `ModuleAiNarrator` /provider runtime; there is no separate provider, prompt, reasoning, or RAG engine to operate.
-

Automation & Incident operations (Sprint 23)

- **Enabling live automation:** live execution is OFF by default. To allow real actions, set `AUTOMATION_LIVE_EXECUTION=true`, enable the specific runner flag(s), and populate `automation.allowed_hosts` for HTTP actions. Re-run config cache. Verify by dispatching a dry-run first, then a live action — confirm it lands in **Approvals** before it runs.
 - **Approving / rolling back:** live actions wait in `GET /api/qynrun/automation/approvals`; an `administerAi` operator decides. To undo a successful, rollback-capable run, call `POST /automation/executions/{uuid}/rollback`.
 - **Triage an incident:** open the QynReact workspace (`GET /api/qynreact/incidents/{uuid}`), review cross-module context, use Copilot/Recommend, mitigate via QynRun, then draft an editable postmortem.
 - **Auditing:** automation events are written via `AutomationAuditLogger` (`action` LIKE `'automation_%'`); incident actions under `incident_*`. Outcomes are captured as auditable learning records — no model training, no hidden state.
 - **Stuck/long runs:** executions honor `automation.default_timeout` and retry policy; check `automation_executions` / `automation_execution_steps` for status and step-level detail.
 - **No duplicated automation:** QynRun/QynReact and all future modules consume the shared Automation Platform — there is no module-specific execution engine to operate. See Docs 24–27.
-

Knowledge, Service Desk & Notification operations (Sprint 24)

- **Check knowledge sources:** `GET /api/qynknow/sources` lists operational vs planned providers. The Internal KB is the only operational source until external providers are wired; planned sources report non-operational and are skipped by Enterprise Search (no fabricated results).
 - **Enterprise Search / Timeline / Graph** are **read-models** over real rows — they write nothing and cannot affect other modules. If results look stale, confirm the underlying rows exist; there is no separate index to rebuild for the Internal KB (it queries `knowledge_documents` directly).
 - **Notification triage:** signals are ingested via `POST /api/qynnotify/notifications` with deterministic deduplication (`dedup_key`), correlation (`correlation_id`), and urgency scoring. Recipients are real members only; if critical signals select no privileged member, the workspace owner is the fallback. Tune `config/knowledge.php → notifications` (severity weights, correlation window) and re-cache config.
 - **Service desk triage:** `POST /api/qynsupport/tickets/{uuid}/intelligence/analyze` returns evidence-based suggestions; when there is no history it reports "**insufficient evidence**" rather than guessing an assignee.
 - **Auditing:** all Sprint 24 events are written via `PlatformAuditLogger` to the shared `audit_logs` table. All data is workspace-isolated and UUID-only.
 - **No duplicated AI/orchestration:** knowledge, ticket, and notification intelligence all narrate through `ModuleAiNarrator` and reuse the `CrossModuleOrchestrator`. See Docs 28–32.
-

Enterprise Intelligence operations (Sprint 25)

- **Platform Health is the first stop:** `GET /api/qynva/health` (requires `administerAi`) reports overall
 - per-area status (AI/automation/knowledge platforms, search, registries, providers, queues, event bus, background jobs). For any `degraded / down` area, read its `reason` and follow the matching registry/ queue/provider check. UI: `/qynva/health`.
- **Event Bus introspection:** `GET /api/qynva/events` (privileged) shows the event vocabulary, registered subscribers (and what each listens to), and the recent event ring. A subscriber that throws is caught and logged — it **never** breaks publishing. Every publish is audited as `platform_event_published`.
- **QynVA never executes:** the operator (`POST /api/qynva/operator/operate`) only returns editable, evidence-based plans referencing existing module actions; humans approve and the owning module executes.
- **QynBalance honesty:** `GET /api/qynbalance/cost/overview` shows real counts and monetary figures **only where** `config/cost.php` **rates are set**; otherwise it reports "pricing unavailable" with a `configure_pricing` recommendation. This is expected behavior, not a fault. Set `COST_*` env to enable monetary estimates and re-cache config.

- **AI disabled:** Executive summaries, QynVA, and the cost copilot fall back to the deterministic mock provider (flagged `ai_enabled:false`); the underlying evidence is always real.
- **Async migration:** Event Bus fan-out is synchronous-but-isolated; replacing the `dispatch()` body with a queued job requires no changes to publishers or subscribers. See Docs 33–36, 44.