



CONFIDENTIAL — PRODUCT

QynAsset Guide (Asset Intelligence)

Qyenx vOPS HUB — v1.0.0 · Document v1.0

Document Version	1.0
Software Version	v1.0.0
Applies To	Qyenx vOPS HUB v1.0.0
Classification	Confidential — Product
Owner	QynAsset Engineering
Status	Released
Last Updated	2026-06-30
Document Type	Module guide

REVISION HISTORY

Version	Date	Notes
1.0	2026-06-30	Initial QynAsset guide: Asset Intelligence (Sprint 22) as the second production AI adapter, grounded strictly in the real discovered inventory.

Table of Contents

- 1. What QynAsset Intelligence is
- 2. Capabilities
- 3. The dashboard
- 4. Contextual AI actions
- 5. Recommendations (always evidence-based)
- 6. API
- 7. Security & governance
- 8. Enabling deeper intelligence

22 — QynAsset Guide (Asset Intelligence)

Audience: Operators, asset managers, and engineers using QynAsset. **Scope:** QynAsset **Asset Intelligence** — the AI surface that explains the real, discovered asset inventory. It is the **second production AI adapter** on the shared Quenyx AI Platform.

1. What QynAsset Intelligence is

QynAsset Intelligence turns the inventory you actually collect into **explainable asset intelligence**. It reuses the shared Quenyx AI Platform end-to-end (provider abstraction, prompt orchestration, conversations, audit) through the **AI Adapter Platform** (see docs 07 and 23) — no AI logic is duplicated.

Ground truth, never invented. An "asset" is a **discovered host**: an `observe_targets_hosts` record enriched by its enrolled agent (`agents`) and that agent's latest inventory push (`agent_inventories`). Hardware/capacity reuses Capacity Planning; dependencies reuse the Infrastructure Map. Capabilities with **no data source** (software licenses; warranty / end-of-life / end-of-support dates) are reported honestly as **not collected** — they are never fabricated.

2. Capabilities

Capability	What it explains	Data source
Asset Discovery Intelligence	New, changed, inactive, unknown, duplicate assets; discovery confidence	hosts + agents + inventory timestamps/identity
CMDB Intelligence	"Which assets are unmanaged / duplicated / without owners?" etc.	inventory summary + discovery
Lifecycle Intelligence	Replacement priority, business impact (warranty/EOL not collected)	agent version/age, activity
Asset Relationship Analysis	Blast radius / SPOF if an asset fails	Infrastructure Map topology
Dependency Intelligence	What an asset depends on / serves; subnet neighbors	service checks + /24 grouping
License Intelligence	Utilization, compliance risk — honest "not collected" until integrated	(none today)
Hardware Intelligence	CPU cores (from inventory) + utilization/growth	inventory + Capacity Planning
Operational Asset Summary	Inventory rollup, with/without agent, online/inactive	inventory summary
Asset Health Summary	Active vs inactive/stale, by OS, by source	agent status + heartbeat
Risk Summary	Evidence-based risks (inactive, duplicate, uncovered, capacity)	recommendations engine

3. The dashboard

`Asset Intelligence` (`/app/workspaces/:id/qynasset/intelligence`) shows real data only: inventory totals (with/without agent, online, inactive), discovery counts (new / changed / unknown / duplicate), inventory by OS, discovery-confidence distribution, inactive and newly-discovered assets with a contextual ✨ **Explain**, evidence-based recommendations, and recent AI investigations. A licensing/lifecycle note states clearly that those facts are not collected and how to enable them.

4. Contextual AI actions

Reusing the Sprint 21 UX (a single sparkle action + the shared Copilot drawer, no duplicated chat):

Context	Action
Asset	✦ Explain
Dependency	✦ Analyze
Lifecycle	✦ Forecast
Relationship	✦ Impact
License	✦ Review

Each opens the **Asset Copilot**, which reuses Quenyx AI conversations — every thread is a real conversation openable in Quenyx AI.

5. Recommendations (always evidence-based)

Examples, each citing concrete evidence: investigate an **inactive** asset (agent offline), **enroll an agent** for a monitored host with no agent (so hardware/inventory becomes available), **review duplicate** assets (same address/name), and plan **hardware/capacity** action from Capacity Planning runway. No recommendation is produced without evidence.

6. API

Workspace-scoped (required `workspace` UUID), UUID-only, behind `throttle:ai-workspace` :

- GET `/api/qynasset/intelligence/overview`
- POST `/api/qynasset/intelligence/copilot`
- GET `/api/qynasset/intelligence/recommendations`
- POST `/api/qynasset/intelligence/assets/{uuid}/explain`
- POST `/api/qynasset/intelligence/assets/{uuid}/dependencies`
- POST `/api/qynasset/intelligence/assets/{uuid}/lifecycle`
- POST `/api/qynasset/intelligence/assets/{uuid}/impact`
- POST `/api/qynasset/intelligence/licenses/review`

See API Reference (doc 08, §19).

7. Security & governance

Every request is workspace-aware, gated by the `qynasset` entitlement, RBAC (`accessAi`), and the `can_use_ai` capability for AI actions; audited, provider-logged, conversation-logged, rate-limited, and UUID-only. No data leakage across workspaces.

8. Enabling deeper intelligence

License and lifecycle-date intelligence require an inventory/license integration (e.g. GLPI/FusionInventory). Until one is connected, those capabilities are exposed but report "not collected"

with the required integration — by design, so the gap is visible rather than faked.