



CONFIDENTIAL — PRODUCT

Incident Response Guide

Quenyx vOPS HUB — v1.0.0 · Document v1.0

Document Version	1.0
Software Version	v1.0.0
Applies To	Quenyx vOPS HUB v1.0.0
Classification	Confidential — Product
Owner	SRE / Operations
Status	Released
Last Updated	2026-06-30
Document Type	Operational playbook

REVISION HISTORY

Version	Date	Notes
1.0	2026-06-30	Initial incident response playbook (Sprint 23) using QynReact + the shared Automation Platform.

Table of Contents

- | 1. Lifecycle
- | 2. Step-by-step
- | 3. Safety rules (non-negotiable)
- | 4. Roles
- | 5. After-action

27 — Incident Response Guide

Audience: On-call engineers and incident commanders. **Scope:** The end-to-end incident lifecycle using QynReact (guide 26), QynRun automation (guide 25), and the shared Automation Platform (guide 24).

1. Lifecycle

```

Detect (alert / manual)
  ↓ open incident (QynReact)
Triage → review cross-module context (QynSight + QynAsset)
  ↓
Investigate → Incident Copilot (grounded Q&A) + Recommend
  ↓
Mitigate → run a runbook / action via QynRun (dry-run → approve → live)
  ↓
Resolve → set resolution, status = resolved
  ↓
Learn → AI-drafted postmortem (editable) + automation learning updated

```

2. Step-by-step

1. **Open** an incident (`POST /api/qynreact/incidents`) with a title and severity. Link an `alert_uuid` / `asset_uuid` when known (deterministic UUIDv5 from QynSight/QynAsset).
2. **Open the workspace** (`GET /incidents/{uuid}`). Review the **Cross-module** section — Operations & Asset Intelligence are reused, never re-collected.
3. **Investigate** with the **Incident Copilot** and **Recommend**. Recommendations cite cross-module evidence and historical automation success rates.
4. **Mitigate** in QynRun:
 - Start with a **dry-run** to preview the plan.
 - For a real fix, dispatch **live** → the action enters **Approvals** → an authorized operator approves → it runs.
 - Use **Rollback** if a successful live action must be undone.
5. **Document** the timeline (`POST /incidents/{uuid}/timeline`) and set the **resolution**.
6. **Postmortem** — generate an editable draft (`POST /incidents/{uuid}/postmortem`), review the root-cause **hypothesis**, finalize action items, and save.

3. Safety rules (non-negotiable)

- No automatic destructive actions — **every live action requires approval**.
- Dry-run is always available and is the default.
- AI never fabricates: missing data (e.g. Knowledge until QynKnow is connected) is reported honestly.
- Everything is workspace-scoped, UUID-only, RBAC-gated, and fully audited.

4. Roles

Action	Minimum role
View incidents / dry-run / copilot	<code>accessAi</code> + <code>can_use_ai</code> (for AI)
Approve live actions / rollback	<code>administerAi</code>

5. After-action

Automation outcomes are captured as auditable learning records; future recommendations cite them. Review the QynRun **Learning** tab periodically to spot low-success actions and refine runbooks.