



QUENYX

V O P S H U B

INTERNAL

Release Notes — Quenyx vOPS HUB v1.0.0

Quenyx vOPS HUB — v1.0.0 · Document v1.0

Document Version	1.0
Software Version	v1.0.0
Applies To	Quenyx vOPS HUB v1.0.0 (GA) — Sprint 25
Classification	Internal
Owner	Platform Engineering
Status	Released
Document Type	Release notes

Table of Contents

Overview

What's new in Sprint 25 (v1.0.0)

Platform Event Bus

Enterprise Context Engine

QynVA — Enterprise AI Operator

QynBalance — Enterprise Cost Intelligence

Executive Intelligence & Enterprise Analytics

Platform Health

Navigation

Compatibility & upgrade

Quality gate (met)

Known limitations

Release Notes — Quenyx vOPS HUB v1.0.0

Overview

Quenyx vOPS HUB v1.0.0 is the first General Availability release and the completion of the original Quenyx roadmap (Phases 1–4, Sprints 1–25). Sprint 25 — the **Enterprise Intelligence Platform** — unifies every module into a single enterprise ecosystem with a shared event bus, a shared context engine, the QynVA Enterprise AI Operator, QynBalance Cost Intelligence, Executive Intelligence, Enterprise Analytics, and Platform Health, and turns on the full module navigation.

What's new in Sprint 25 (v1.0.0)

Platform Event Bus

Shared publish/subscribe layer (`PlatformEventBus`) with a 21-event canonical vocabulary (`PlatformEventNames`), immutable `PlatformEvent` DTO, `EventSubscriber` contract, audited publishes, async-ready fan-out, and an example `NotificationFanoutSubscriber` . Eliminates direct module-to-module calls.

Enterprise Context Engine

A single `EnterpriseContextEngine` that assembles one normalized AI context from workspace, user, permissions, monitoring, assets, automation, knowledge, incidents, notifications, compliance, timeline, graph and search — reusing existing read-models (cross-module gather, Global Timeline, Knowledge Graph v2, Enterprise Search). Recursion-safe; consumed by all AI surfaces.

QynVA — Enterprise AI Operator

Not a chatbot: discovers adapters/capabilities, builds enterprise context, reasons through `ModuleAiNarrator` , and proposes **editable, evidence-based** cross-module coordination plans. Never executes; every turn is recorded and publishes `ConversationCompleted` . New `/api/qynva/*` surface.

QynBalance — Enterprise Cost Intelligence

Cost analysis over **real** platform data (hosts, services, agents, seats, automation). Capacity/license/ cloud optimization, automation savings, asset utilization, budget forecasting, and advisory recommendations. **No fabricated financial values** — monetary figures appear only where real rates are configured; otherwise counts + "pricing unavailable". New `/api/qynbalance/*` surface.

Executive Intelligence & Enterprise Analytics

Evidence-based executive dashboard (operational/infrastructure/compliance health, capacity forecast, top risks/recommendations, automation success, AI usage, incident/knowledge/cost KPIs) with an AI executive summary; and the shared analytics platform (MTTD, MTTR, trends, effectiveness, adoption, KPIs). Honest `available: false` when data is insufficient.

Platform Health

Self-monitoring of AI/automation/knowledge platforms, search, registries, provider health, queues, event bus, and background jobs, with overall + per-area status.

Navigation

The temporary sidebar feature flag is removed. **QynSight, QynAsset, QynRun, QynReact, QynKnow, QynSupport, QynNotify, QynShield, QynBalance, QynVA** are all enabled. Platform items (Dashboard, Workspaces, Quenyx AI, Integrations) retained. **QynCore** remains platform-only.

Compatibility & upgrade

- **No breaking changes** to Sprint 20–24 behavior. All previous APIs, routes, and data are unchanged.
- New optional config: `config/cost.php` (+ `COST_*` env) for QynBalance pricing. Without it, QynBalance is fully functional in count/"pricing unavailable" mode.
- See the **Migration Guide** (doc 40) for the exact upgrade steps.

Quality gate (met)

No duplicated platform logic · event-driven · shared Context Engine / AI / Automation / Knowledge · workspace isolation · UUID-only · RBAC · audit · EN/AR complete · backend syntax clean · documentation updated · PDFs regenerated.

Known limitations

- AI narration runs on the deterministic mock provider unless an AI provider is configured (evidence is always real; responses are flagged `ai_enabled: false`).
- QynBalance monetary estimates require operator-supplied rates; cloud-spend optimization needs a connected cloud billing source.
- Event Bus fan-out is synchronous-but-isolated; the queue-backed path is a drop-in (no publisher/subscriber changes) but ships disabled by default.